Processor FSM Design

Mainak Chaudhuri
Indian Institute of Technology Kanpur

Sketch

- · Abstract model of computer
- Single-cycle instruction execution
- Multi-cycle instruction execution
- Pipelined instruction execution

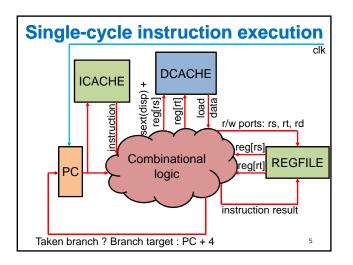
2

Abstract model of computer

- Each instruction undergoes five stages
 - Stage 0 (IF): fetch the instruction pointed to by program counter from memory
 - Stage 1 (ID/RF): decode the instruction to extract various fields and read source register operands
 - Source registers can be read in parallel with instruction decoding in MIPS due to fixed positions of rs and rt in all formats of the ISA
 - Stage 2 (EX): execute the instruction in ALU;
 compute address of load/store instructions;
 update program counter to PC+4 or branch
 target (if this instruction is control transfer ins)
 - Stage 3 (MEM): access memory if load/store instruction; use address computed in stage 2
 - Stage 4 (WB): write result back to destination register if the instruction produces a result

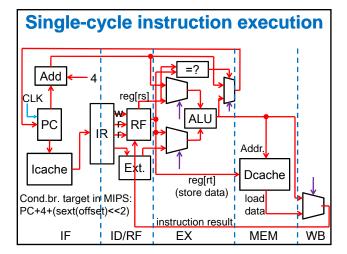
Single-cycle instruction execution

- Possible to implement all five stages as a big combinational logic
 - All the memory elements are also combinational meaning that the input address can be presented any time within a cycle and the output data is available within the same cycle
 - Register file (in stages 1 and 4), icache (in stage 0), and dcache (in stage 3) are the combinational memory elements
 - PC, icache[PC], reg[rs], reg[rt], dcache[disp(rs)] are inputs
 dcache[disp(rs)] is a relevant input for loads
 - PC, reg[rd] or reg[rt], dcache[disp(rs)] are outputs
 dcache[disp(rs)] is a relevant output for stores
 - The instruction opcode and function (for R format) fields decide the control of combinational logic e.g., which register operand(s) is/are valid sources, which ALU operation to invoke, etc.



Single-cycle instruction execution

- Taken branch
 - A branch instruction whose condition dictates that the next instruction to be executed is the one at the branch target
 - beg \$1, \$2, label (the branch is taken if \$1 == \$2)
 - Unconditional direct and indirect jumps and procedure calls are all taken branches
- · Not taken branch
 - A branch instruction whose condition dictates that the next instruction to be executed is the instruction right next to the branch
 - A la Frost's "Road Not Taken"
 - beg \$1, \$2, label (the branch is not taken if \$1 != \$2)
 - A conditional branch can be taken or not taken depending on how the condition evaluates at run time



Single-cycle instruction execution

- Clock cycle time must be large enough to accommodate the lengthiest instruction
 - This is typically the load instruction whose critical path includes all five stages: IF: fetch instruction, ID/RF: decode while reading reg[rs] and sign extending displacement, EX: add reg[rs] with sext(disp) to generate address, MEM: read dcache to get load data, WB: write load data back to reg[rt]
 - Not a good idea since many other instructions take smaller amount of time
 - Control transfer instructions require only three stages
 - Stores require only four stages
 - · All ALU instructions can bypass MEM stage

Multi-cycle instruction execution

- Each stage takes one cycle to execute
 - Processor is now a five state FSM and the memory elements are sequential
 - Need flip-flops at stage boundaries and each stage is a combinational logic
 - Clock cycle time = latency of the longest stage
 - Most instructions take five cycles
 - ALU instructions do nothing in MEM stage, but will have to go through it if FSM state transition is implemented using an incrementer
 - Some instructions may take longer e.g., multiply/divide, load/store (in case of dcache miss)
 - Some instructions may take shorter time e.g., stores take four cycles, control transfer instructions take three cycles
 - Overall CPI depends on instruction mix

Multi-cycle instruction execution

- Example
 - IF: 2 ns, ID/RF: 1 ns, EX: 1 ns, DMEM: 3 ns, WB: 1 ns
 - Branch frequency: 20%
 - Store frequency: 10%
 - Multiply/divide frequency: 5%, latency: 30 ns
 - Total instruction count: 100
- Multi-cycle
 - Cycle time: 3 ns, frequency: 333 MHz
 - CPI: 0.2*3+0.1*4+0.05*10+0.65*5 = 4.75
 - Execution time: 100*4.75*3 ns = 1425 ns
- Single-cycle
 - Cycle time: 8 ns, frequency: 125 MHz
 - CPI: 1*0.95 + ceil(30/8)*0.05 = 1.15
 - Execution time: 100*1.15*8 ns = 920 ns

10

Multi-cycle instruction execution

- Same example with a balanced pipeline
 - IF: 2 ns, ID/RF: 2 ns, EX: 2 ns, DMEM: 2 ns, WB: 2 ns
 - Branch frequency: 20%
 - Store frequency: 10%
 - Multiply/divide frequency: 5%, latency: 30 ns
 - Total instruction count: 100
- Multi-cycle
 - Cycle time: 2 ns, frequency: 500 MHz
 - CPI: 0.2*3+0.1*4+0.05*15+0.65*5 = 5
 - Execution time: 100*5*2 ns = 1000 ns
- Single-cycle
 - Cycle time: 10 ns, frequency: 100 MHz
 - CPI: 1*0.95 + ceil(30/10)*0.05 = 1.1
 - Execution time: 100*1.1*10 ns = 1100 ns (worse than multi-cycle)

Pipelined instruction execution

- Observations from multi-cycle design
 - In the second cycle, I know if it is a branch; if not, start fetching the next instruction?
 - When the ALU is doing an addition (say), the decoder is sitting idle; can we use it for some other instruction?
 - In summary, exactly one stage is active at any point in time: wastes hardware resources
- Form a pipeline
 - Process five instructions in parallel
 - Each instruction is in a different stage of processing (called pipe stage)

Pipelined instruction execution

- Individual instruction latency is five cycles, but ideally can finish one instruction every cycle after the pipeline is filled up
 - Ideal CPI of 1.0 at the clock frequency of multicycle design
 - Execution time is ideally one-fifth of the multicycle design
 - Instruction throughput improves five times (number of instructions completed in a given time)

10	IF	ID/RF	EX	MEM	WB			
11		IF	ID/RF	EX	MEM	WB		
12			IF	ID/RF	EX	MEM	WB	
13				IF	ID/RF	EX	MEM	WB
14			tim o		IF	ID/RF	EX	M€M
		$\overline{}$	IIIIie —					

Pipelined instruction execution

- Gains
 - Extracting parallelism from a sequential instruction stream: known as instruction-level parallelism (ILP)
 - Can complete one instruction every cycle (ideally)
- Loss
 - Each pipe stage may get lengthened a little bit due to control overhead (skew time, setup time, propagation delay): limits the gain due to pipelining
 - Each instruction may take slightly longer for this reason
 - Bigger aggregate memory bandwidth: icache and dcache may miss in the same cycle
 - Pipeline hazards

. . .

Pipeline hazard: control

• Branches pose a problem

Cycles 7

ins *i* (*beq* \$1, \$0, label) IF ID/RF EX MEM WB What to fetch? IF What to fetch?

ins i+1 (target)

- Two pipeline bubbles: increases average CPI
- Can we reduce it to one bubble?

15

Branch delay slot

- MIPS R3000 has one bubble
 - Called branch delay slot
 - Exploit clock cycle phases
 - On the positive half compute branch condition
 - On the negative half fetch the target

beq \$1, \$0,label IF ID EX MEM WB Bubble

Target

 The PC update hardware (selection between target and next PC) works on the lower edge

Branch delay slot

- · Can we utilize the branch delay slot?
 - The delay slot is always executed (irrespective of the fate of the branch)
 - Reason why branch target is PC+4+(sext(offset)<<2)
 - Boost instructions common to fall through and target paths to the delay slot or from earlier than the branch
 - Not always possible to find
 - Must boost something that does not alter the outcome of the computation
 - If the BD slot is filled with useful instruction then we don't lose anything in CPI; otherwise we pay a branch penalty of one cycle

Branch prediction

- Today all processors rely on branch predictors that observe the behavior of individual branches and learn to predict their future behavior
 - Makes it possible to infer the next instruction's PC even before the branch executes
 - Pipeline must be flushed on a wrong prediction

18

Pipeline hazard: data

- Pipelining disturbs the sequential thoughtprocess
 - Data dependencies among instructions start to show up

add \$1, \$2, \$3 sub \$4, \$1, \$5 and \$6, \$1, \$7 or \$8, \$1, \$9 xor \$10, \$1, \$11

- Result of add is needed by all instructions (RAW hazard: read after write hazard) IF ID EX MEM WB add IF ID EX MEM WB sub IF ID EX MEM WB and IF ID EX MEM WB or xor IF ď EX MEM WB

Pipeline hazard: data

- How to avoid increasing CPI?
 - Can we forward the correct value just in time?

add IF ID EX MEM WB sub IF ID EX MEM WB and IF ID EX MEM WB

 Read wrong value in ID/RF, but bypassed value overrides it (need a multiplexor for each ALU input to choose between the RF value and the bypassed value)

Pipeline hazard: data

• Can we always avoid stalling?

lw \$1, 0(\$2) sub \$4, \$1, \$5 and \$6, \$1, \$7 or \$8, \$1, \$9

Iw IF ID EX MEM WB sub IF ID EX MEM WB and IF ID EX MEM WB or IF ID EX MEM WB

- Need some time travel (backwards)! Not yet feasible!!
- One option: hardware pipeline interlock to stall the sub by a cycle
- Early generations of MIPS (Microprocessor without Interlocked Pipe Stages) had the compiler to fill the load delay slot with something independent or a NOP

Pipelined instruction execution

- Hazards can cause pipeline stalls and introduce bubbles in the pipeline depending on the pipeline organization
- Overall speedup of pipelining over multicycle non-pipelined implementation = number of pipe stages/(1 + average stall cycles per instruction)